

10 minutes to pandas

This is a short introduction to pandas, geared mainly for new users. You can see more complex recipes in the [Cookbook](#).

Customarily, we import as follows:

```
In [1]: import numpy as np
In [2]: import pandas as pd
```

Object creation

See the [Intro to data structures section](#).

Creating a **Series** by passing a list of values, letting pandas create a default integer index:

```
In [3]: s = pd.Series([1, 3, 5, np.nan, 6, 8])
In [4]: s
Out[4]:
0    1.0
1    3.0
2    5.0
3    NaN
4    6.0
5    8.0
dtype: float64
```

Creating a **DataFrame** by passing a NumPy array, with a datetime index and labeled columns:

```
In [5]: dates = pd.date_range("20130101", periods=6)
In [6]: dates
Out[6]:
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')
In [7]: df = pd.DataFrame(np.random.randn(6, 4), index=dates,
                           columns=list("ABCD"))
In [8]: df
Out[8]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

Creating a **DataFrame** by passing a dictionary of objects that can be converted into a series-like structure:

```
In [9]: df2 = pd.DataFrame(
...:     {
...:         "A": 1.0,
...:         "B": pd.Timestamp("20130102"),
...:         "C": pd.Series(1, index=list(range(4)), dtype="float32"),
...:         "D": np.array([3] * 4, dtype="int32"),
...:         "E": pd.Categorical(["test", "train", "test", "train"]),
...:         "F": "foo",
...:     }
...: )
```

```
In [10]: df2
Out[10]:
```

	A	B	C	D	E	F
0	1.0	2013-01-02	1.0	3	test	foo
1	1.0	2013-01-02	1.0	3	train	foo
2	1.0	2013-01-02	1.0	3	test	foo
3	1.0	2013-01-02	1.0	3	train	foo

Series = combo of 2 arrays

The columns of the resulting DataFrame have different dtypes:

```
In [11]: df2.dtypes
Out[11]:
A      float64
B  datetime64[ns]
C      float32
D       int32
E      category
F      object
dtype: object
```

If you're using IPython, tab completion for column names (as well as public attributes) is automatically enabled. Here's a subset of the attributes that will be completed:

```
In [12]: df2.<TAB>
df2.A      df2.bool
df2.abs     df2.boxplot
df2.add     df2.C
df2.add_prefix df2.clip
df2.add_suffix df2.columns
df2.align   df2.copy
df2.all     df2.count
df2.any     df2.combine
df2.append  df2.D
df2.apply   df2.describe
df2.applymap df2.diff
df2.B       df2.duplicated
```

As you can see, the columns A, B, C, and D are automatically tab completed. E and F are there as well; the rest of the attributes have been truncated for brevity.

Viewing data

See the [Basics section](#).

Here is how to view the top and bottom rows of the frame:

```
In [13]: df.head()
Out[13]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

```
In [14]: df.tail(3)
Out[14]:
```

	A	B	C	D
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

Display the index, columns:

```
In [15]: df.index
Out[15]:
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')

In [16]: df.columns
Out[16]: Index(['A', 'B', 'C', 'D'], dtype='object')
```

`DataFrame.to_numpy()` gives a NumPy representation of the underlying data. Note that this can be an expensive operation when your `DataFrame` has columns with different data types, which comes down to a fundamental difference between pandas and NumPy: **NumPy arrays have one dtype for the entire array, while pandas DataFrames have one dtype per column.** When you call `DataFrame.to_numpy()`, pandas will find the NumPy dtype that can hold *all* of the dtypes in the `DataFrame`. This may end up being `object`, which requires casting every value to a Python object.

For `df`, our `DataFrame` of all floating-point values, `DataFrame.to_numpy()` is fast and doesn't require copying data:

```
In [17]: df.to_numpy()
Out[17]:
array([[ 0.4691, -0.2829, -1.5091, -1.1356],
       [ 1.2121, -0.1732,  0.1192, -1.0442],
       [-0.8618, -2.1046, -0.4949,  1.0718],
       [ 0.7216, -0.7068, -1.0396,  0.2719],
       [-0.425 ,  0.567 ,  0.2762, -1.0874],
       [-0.6737,  0.1136, -1.4784,  0.525 ]])
```

For `df2`, the `DataFrame` with multiple dtypes, `DataFrame.to_numpy()` is relatively expensive

```
In [18]: df2.to_numpy()
Out[18]:
array([[1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'test', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'train', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'test', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'train', 'foo']],
      dtype=object)
```

Note

`DataFrame.to_numpy()` does *not* include the index or column labels in the output.

`describe()` shows a quick statistic summary of your data:

```
In [19]: df.describe()
Out[19]:
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	0.073711	-0.431125	-0.687758	-0.233103
std	0.843157	0.922818	0.779887	0.973118
min	-0.861849	-2.104569	-1.509059	-1.135632
25%	-0.611510	-0.600794	-1.368714	-1.076610
50%	0.022070	-0.228039	-0.767252	-0.386188
75%	0.658444	0.041933	-0.034326	0.461706
max	1.212112	0.567020	0.276232	1.071804

Transposing your data:

```
In [20]: df.T
Out[20]:
```

	2013-01-01	2013-01-02	2013-01-03	2013-01-04	2013-01-05	2013-01-06
A	0.469112	1.212112	-0.861849	0.721555	-0.424972	-0.673690
B	-0.282863	-0.173215	-2.104569	-0.706771	0.567020	0.113648
C	-1.509059	0.119209	-0.494929	-1.039575	0.276232	-1.478427
D	-1.135632	-1.044236	1.071804	0.271860	-1.087401	0.524988

Sorting by an axis:

(column header) desc sort

```
In [21]: df.sort_index(axis=1, ascending=False)
```

```
Out[21]:
```

	D	C	B	A
2013-01-01	-1.135632	-1.509059	-0.282863	0.469112
2013-01-02	-1.044236	0.119209	-0.173215	1.212112
2013-01-03	1.071804	-0.494929	-2.104569	-0.861849
2013-01-04	0.271860	-1.039575	-0.706771	0.721555
2013-01-05	-1.087401	0.276232	0.567020	-0.424972
2013-01-06	0.524988	-1.478427	0.113648	-0.673690

Sorting by values:

```
In [22]: df.sort_values(by="B")
```

```
Out[22]:
```

	A	B	C	D
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-06	-0.673690	0.113648	-1.478427	0.524988
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

Selection

Note

While standard Python / NumPy expressions for selecting and setting are intuitive and come in handy for interactive work, for production code, we recommend the optimized pandas data access methods, `.at`, `.iat`, `.loc` and `.iloc`.

See the [indexing documentation](#) [Indexing and Selecting Data](#) and [MultiIndex / Advanced Indexing](#).

Getting

Selecting a single column, which yields a `Series` equivalent to `df.A`:

```
In [23]: df["A"]
```

```
Out[23]:
```

2013-01-01	0.469112
2013-01-02	1.212112
2013-01-03	-0.861849
2013-01-04	0.721555
2013-01-05	-0.424972
2013-01-06	-0.673690

Freq: D, Name: A, dtype: float64

Selecting via `[],` which slices the rows:

```
In [24]: df[0:3]
```

```
Out[24]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804

```
In [25]: df["20130102":"20130104"]
```

```
Out[25]:
```

	A	B	C	D
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860

Selection by label

See more in [Selection by Label](#).

For getting a cross section using a label:

```
In [26]: df.loc[dates[0]]
Out[26]:
A    0.469112
B   -0.282863
C   -1.509059
D   -1.135632
Name: 2013-01-01 00:00:00, dtype: float64
```

1st column

Selecting on a multi-axis by label:

```
In [27]: df.loc[:, ["A", "B"]]
Out[27]:
```

	A	B
2013-01-01	0.469112	-0.282863
2013-01-02	1.212112	-0.173215
2013-01-03	-0.861849	-2.104569
2013-01-04	0.721555	-0.706771
2013-01-05	-0.424972	0.567020
2013-01-06	-0.673690	0.113648

Showing label slicing, both endpoints are included:

```
In [28]: df.loc["20130102":"20130104", ["A", "B"]]
Out[28]:
```

	A	B
2013-01-02	1.212112	-0.173215
2013-01-03	-0.861849	-2.104569
2013-01-04	0.721555	-0.706771

Reduction in the dimensions of the returned object:

```
In [29]: df.loc["20130102", ["A", "B"]]
Out[29]:
A    1.212112
B   -0.173215
Name: 2013-01-02 00:00:00, dtype: float64
```

For getting a scalar value:

```
In [30]: df.loc[dates[0], "A"]
Out[30]: 0.4691122999071863
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [31]: df.at[dates[0], "A"]
Out[31]: 0.4691122999071863
```

Selection by position

See more in [Selection by Position](#).

Select via the position of the passed integers:

```
In [32]: df.iloc[3]
Out[32]:
A    0.721555
B   -0.706771
C   -1.039575
D    0.271860
Name: 2013-01-04 00:00:00, dtype: float64
```

By integer slices, acting similar to NumPy/Python:

stop value excluded

```
In [33]: df.iloc[3:5, 0:2]
Out[33]:
```

	A	B
2013-01-04	0.721555	-0.706771
2013-01-05	-0.424972	0.567020

(3-4)

By lists of integer position locations, similar to the NumPy/Python style:

not range

```
In [34]: df.iloc[[1, 2, 4], [0, 2]]
Out[34]:
```

	A	C
2013-01-02	1.212112	0.119209
2013-01-03	-0.861849	-0.494929
2013-01-05	-0.424972	0.276232

0 1 2 4, *→ list '0'*, *→ stop range*

For slicing rows explicitly:

1:3, *all column*

```
In [35]: df.iloc[1:3, :]
Out[35]:
```

	A	B	C	D
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804

1:3, *excluded*

For slicing columns explicitly:

all rows, *2 col*

```
In [36]: df.iloc[:, 1:3]
Out[36]:
```

	B	C
2013-01-01	-0.282863	-1.509059
2013-01-02	-0.173215	0.119209
2013-01-03	-2.104569	-0.494929
2013-01-04	-0.706771	-1.039575
2013-01-05	0.567020	0.276232
2013-01-06	0.113648	-1.478427

For getting a value explicitly:

```
In [37]: df.iloc[1, 1]
Out[37]: -0.17321464905330858
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [38]: df.iat[1, 1]
Out[38]: -0.17321464905330858
```

Boolean indexing

Using a single column's values to select data:

only A the values

```
In [39]: df[df["A"] > 0]
Out[39]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-04	0.721555	-0.706771	-1.039575	0.271860

Selecting values from a DataFrame where a boolean condition is met:

```
In [40]: df[df > 0]
Out[40]:
```

	A	B	C	D
2013-01-01	0.469112	NaN	NaN	NaN
2013-01-02	1.212112	NaN	0.119209	NaN
2013-01-03	NaN	NaN	NaN	1.071804
2013-01-04	0.721555	NaN	NaN	0.271860
2013-01-05	NaN	0.567020	0.276232	NaN
2013-01-06	NaN	0.113648	NaN	0.524988

Using the `isin()` method for filtering:

```
In [41]: df2 = df.copy()
In [42]: df2["E"] = ["one", "one", "two", "three", "four", "three"]
In [43]: df2
Out[43]:
```

	A	B	C	D	E
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632	one
2013-01-02	1.212112	-0.173215	0.119209	-1.044236	one
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804	two
2013-01-04	0.721555	-0.706771	-1.039575	0.271860	three
2013-01-05	-0.424972	0.567020	0.276232	-1.087401	four
2013-01-06	-0.673690	0.113648	-1.478427	0.524988	three

f

```
In [44]: df2[df2["E"].isin(["two", "four"])]
Out[44]:
```

	A	B	C	D	E
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804	two
2013-01-05	-0.424972	0.567020	0.276232	-1.087401	four

Setting

Setting a new column automatically aligns the data by the indexes:

```
In [45]: s1 = pd.Series([1, 2, 3, 4, 5, 6], index=pd.date_range("20130102",
periods=6))
In [46]: s1
Out[46]:
```

2013-01-02	1
2013-01-03	2
2013-01-04	3
2013-01-05	4
2013-01-06	5
2013-01-07	6

Freq: D, dtype: int64

```
In [47]: df["F"] = s1
```

Setting values by label:

```
In [48]: df.at[dates[0], "A"] = 0
```

Setting values by position:

```
In [49]: df.iat[0, 1] = 0
```

Setting by assigning with a NumPy array:

```
In [50]: df.loc[:, "D"] = np.array([5] * len(df))
```

The result of the prior setting operations:

```
In [51]: df
Out[51]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-1.509059	5	NaN
2013-01-02	1.212112	-0.173215	0.119209	5	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5	2.0
2013-01-04	0.721555	-0.706771	-1.039575	5	3.0
2013-01-05	-0.424972	0.567020	0.276232	5	4.0
2013-01-06	-0.673690	0.113648	-1.478427	5	5.0

A **where** operation with **setting**:

```
In [52]: df2 = df.copy()
In [53]: df2[df2 > 0] = -df2
```

all negative

```
In [54]: df2
Out[54]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-1.509059	-5	NaN
2013-01-02	-1.212112	-0.173215	-0.119209	-5	-1.0
2013-01-03	-0.861849	-2.104569	-0.494929	-5	-2.0
2013-01-04	-0.721555	-0.706771	-1.039575	-5	-3.0
2013-01-05	-0.424972	-0.567020	-0.276232	-5	-4.0
2013-01-06	-0.673690	-0.113648	-1.478427	-5	-5.0

Missing data

pandas primarily uses the value `np.nan` to represent missing data. It is by default not included in computations. See the [Missing Data section](#).

Reindexing allows you to change/add/delete the index on a specified axis. This returns a copy of the data:

```
In [55]: df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
In [56]: df1.loc[dates[0] : dates[1], "E"] = 1
In [57]: df1
Out[57]:
```

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-1.509059	5	NaN	1.0
2013-01-02	1.212112	-0.173215	0.119209	5	1.0	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5	2.0	NaN
2013-01-04	0.721555	-0.706771	-1.039575	5	3.0	NaN

To drop any rows that have missing data:

```
In [58]: df1.dropna(how="any")
Out[58]:
```

	A	B	C	D	F	E
2013-01-02	1.212112	-0.173215	0.119209	5	1.0	1.0

Filling missing data:

```
In [59]: df1.fillna(value=5)
Out[59]:
```

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-1.509059	5	5.0	1.0
2013-01-02	1.212112	-0.173215	0.119209	5	1.0	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5	2.0	5.0
2013-01-04	0.721555	-0.706771	-1.039575	5	3.0	5.0

To get the boolean mask where values are nan:

```
In [60]: pd.isna(df1)
Out[60]:
```

	A	B	C	D	F	E
2013-01-01	False	False	False	False	True	False
2013-01-02	False	False	False	False	False	False
2013-01-03	False	False	False	False	False	True
2013-01-04	False	False	False	False	False	True

nan (circled around F in row 1 and E in row 4)

Operations

See the [Basic section on Binary Ops](#).

Stats

Operations in general *exclude* missing data.

Performing a descriptive statistic:

```
In [61]: df.mean()  
Out[61]: columnwise  
A    -0.004474  
B    -0.383981  
C    -0.687758  
D     5.000000  
F     3.000000  
dtype: float64
```

Same operation on the other axis:

```
In [62]: df.mean(1) Rowwise  
Out[62]:  
2013-01-01    0.872735  
2013-01-02    1.431621  
2013-01-03    0.707731  
2013-01-04    1.395042  
2013-01-05    1.883656  
2013-01-06    1.592306  
Freq: D, dtype: float64
```

Operating with objects that have *different dimensionality* and *need alignment*. In addition, pandas automatically broadcasts along the specified dimension:

```
In [63]: s = pd.Series([1, 3, 5, np.nan, 6, 8], index=dates).shift(2)  
  
In [64]: s  
Out[64]:  
2013-01-01    NaN  
2013-01-02    NaN  
2013-01-03    1.0  
2013-01-04    3.0  
2013-01-05    5.0  
2013-01-06    NaN  
Freq: D, dtype: float64  
  
In [65]: df.sub(s, axis="index")  
Out[65]:
```

	A	B	C	D	F
2013-01-01	NaN	NaN	NaN	NaN	NaN
2013-01-02	NaN	NaN	NaN	NaN	NaN
2013-01-03	-1.861849	-3.104569	-1.494929	4.0	1.0
2013-01-04	-2.278445	-3.706771	-4.039575	2.0	0.0
2013-01-05	-5.424972	-4.432980	-4.723768	0.0	-1.0
2013-01-06	NaN	NaN	NaN	NaN	NaN

Apply

Applying functions to the data:

```
In [66]: df.apply(np.cumsum) ?  
Out[66]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-1.509059	5	NaN
2013-01-02	1.212112	-0.173215	-1.389850	10	1.0
2013-01-03	0.350263	-2.277784	-1.884779	15	3.0
2013-01-04	1.071818	-2.984555	-2.924354	20	6.0
2013-01-05	0.646846	-2.417535	-2.648122	25	10.0
2013-01-06	-0.026844	-2.303886	-4.126549	30	15.0

```
In [67]: df.apply(lambda x: x.max() - x.min())  
Out[67]:  
A    2.073961  
B    2.671590  
C    1.785291  
D    0.000000  
F    4.000000  
dtype: float64
```

Histogramming

See more at [Histogramming and Discretization](#).

```
In [68]: s = pd.Series(np.random.randint(0, 7, size=10))

In [69]: s
Out[69]:
0    4
1    2
2    1
3    2
4    6
5    4
6    4
7    6
8    4
9    4
dtype: int64

In [70]: s.value_counts()
Out[70]:
4    5
2    2
6    2
1    1
dtype: int64
```

→ Histogram for job

String Methods

`Series` is equipped with a set of string processing methods in the `str` attribute that make it easy to operate on each element of the array, as in the code snippet below. Note that pattern-matching in `str` generally uses regular expressions by default (and in some cases always uses them). See more at [Vectorized String Methods](#).

```
In [71]: s = pd.Series(["A", "B", "C", "Aaba", "Baca", np.nan, "CABA", "dog", "cat"])

In [72]: s.str.lower()
Out[72]:
0    a
1    b
2    c
3    aaba
4    baca
5    NaN
6    caba
7    dog
8    cat
dtype: object
```

Merge

Concat

pandas provides various facilities for easily combining together `Series` and `DataFrame` objects with various kinds of set logic for the indexes and relational algebra functionality in the case of join / merge-type operations.

See the [Merging section](#).

Concatenating pandas objects together with `concat()`.

```
In [73]: df = pd.DataFrame(np.random.randn(10, 4))
```

```
In [74]: df
```

```
Out[74]:
```

	0	1	2	3
0	-0.548702	1.467327	-1.015962	-0.483075
1	1.637550	-1.217659	-0.291519	-1.745505
2	-0.263952	0.991460	-0.919069	0.266046
3	-0.709661	1.669052	1.037882	-1.705775
4	-0.919854	-0.042379	1.247642	-0.009920
5	0.290213	0.495767	0.362949	1.548106
6	-1.131345	-0.089329	0.337863	-0.945867
7	-0.932132	1.956030	0.017587	-0.016692
8	-0.575247	0.254161	-1.143704	0.215897
9	1.193555	-0.077118	-0.408530	-0.862495

```
# break it into pieces
```

```
In [75]: pieces = [df[:3], df[3:7], df[7:]]
```

```
In [76]: pd.concat(pieces)
```

```
Out[76]:
```

	0	1	2	3
0	-0.548702	1.467327	-1.015962	-0.483075
1	1.637550	-1.217659	-0.291519	-1.745505
2	-0.263952	0.991460	-0.919069	0.266046
3	-0.709661	1.669052	1.037882	-1.705775
4	-0.919854	-0.042379	1.247642	-0.009920
5	0.290213	0.495767	0.362949	1.548106
6	-1.131345	-0.089329	0.337863	-0.945867
7	-0.932132	1.956030	0.017587	-0.016692
8	-0.575247	0.254161	-1.143704	0.215897
9	1.193555	-0.077118	-0.408530	-0.862495

Note

Adding a column to a `DataFrame` is relatively fast. However, adding a row requires a copy, and may be expensive. We recommend passing a pre-built list of records to the `DataFrame` constructor instead of building a `DataFrame` by iteratively appending records to it.

Join

SQL style merges. See the [Database style joining](#) section.

```
In [77]: left = pd.DataFrame({"key": ["foo", "foo"], "lval": [1, 2]})
```

```
In [78]: right = pd.DataFrame({"key": ["foo", "foo"], "rval": [4, 5]})
```

```
In [79]: left
```

```
Out[79]:
```

	key	lval
0	foo	1
1	foo	2

```
In [80]: right
```

```
Out[80]:
```

	key	rval
0	foo	4
1	foo	5

```
In [81]: pd.merge(left, right, on="key")
```

```
Out[81]:
```

	key	lval	rval
0	foo	1	4
1	foo	1	5
2	foo	2	4
3	foo	2	5

Another example that can be given is:

```

In [82]: left = pd.DataFrame({"key": ["foo", "bar"], "lval": [1, 2]})
In [83]: right = pd.DataFrame({"key": ["foo", "bar"], "rval": [4, 5]})
In [84]: left
Out[84]:
  key  lval
0  foo     1
1  bar     2
In [85]: right
Out[85]:
  key  rval
0  foo     4
1  bar     5
In [86]: pd.merge(left, right, on="key")
Out[86]:
  key  lval  rval
0  foo     1     4
1  bar     2     5

```

Grouping

By "group by" we are referring to a process involving one or more of the following steps:

- **Splitting** the data into groups based on some criteria
- **Applying** a function to each group independently
- **Combining** the results into a data structure

See the [Grouping section](#).

```

In [87]: df = pd.DataFrame(
...:     {
...:         "A": ["foo", "bar", "foo", "bar", "foo", "bar", "foo", "foo"],
...:         "B": ["one", "one", "two", "three", "two", "two", "one", "one"],
...:         "C": np.random.randn(8),
...:         "D": np.random.randn(8),
...:     }
...: )
In [88]: df
Out[88]:
   A      B      C      D
0  foo  one  1.346061 -1.577585
1  bar  one  1.511763  0.396823
2  foo  two  1.627081 -0.105381
3  bar three -0.990582 -0.532532
4  foo  two -0.441652  1.453749
5  bar  two  1.211526  1.208843
6  foo  one  0.268520 -0.080952
7  foo three  0.024580 -0.264610

```

Grouping and then applying the `sum()` function to the resulting groups:

```

In [89]: df.groupby("A").sum()
Out[89]:
      C      D
A
bar  1.732707  1.073134
foo  2.824590 -0.574779

```

Grouping by multiple columns forms a hierarchical index, and again we can apply the `sum()` function:

```
In [90]: df.groupby(["A", "B"]).sum()
```

```
Out[90]:
```

		C	D
bar	one	1.511763	0.396823
	three	-0.990582	-0.532532
	two	1.211526	1.208843
foo	one	1.614581	-1.658537
	three	0.024580	-0.264610
	two	1.185429	1.348368

Reshaping

See the sections on [Hierarchical Indexing](#) and [Reshaping](#).

Stack

```
In [91]: tuples = list(
...:     zip(
...:         *[
...:             ["bar", "bar", "baz", "baz", "foo", "foo", "qux", "qux"],
...:             ["one", "two", "one", "two", "one", "two", "one", "two"],
...:         ]
...:     )
...: )
```

```
In [92]: index = pd.MultiIndex.from_tuples(tuples, names=["first", "second"])
```

```
In [93]: df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=["A", "B"])
```

```
In [94]: df2 = df[:4]
```

```
In [95]: df2
```

```
Out[95]:
```

		A	B
bar	one	-0.727965	-0.589346
	two	0.339969	-0.693205
baz	one	-0.339355	0.593616
	two	0.884345	1.591431

The `stack()` method “compresses” a level in the DataFrame’s columns:

```
In [96]: stacked = df2.stack()
```

```
In [97]: stacked
```

```
Out[97]:
```

first	second	
bar	one	A -0.727965
		B -0.589346
	two	A 0.339969
		B -0.693205
baz	one	A -0.339355
		B 0.593616
	two	A 0.884345
		B 1.591431

dtype: float64

With a “stacked” DataFrame or Series (having a `MultiIndex` as the `index`), the inverse operation of `stack()` is `unstack()`, which by default unstacks the **last level**:

```

In [98]: stacked.unstack()
Out[98]:
           A      B
first second
bar  one  -0.727965 -0.589346
     two   0.339969 -0.693205
baz   one  -0.339355  0.593616
     two   0.884345  1.591431

In [99]: stacked.unstack(1)
Out[99]:
second      one      two
first
bar  A -0.727965  0.339969
     B -0.589346 -0.693205
baz  A -0.339355  0.884345
     B  0.593616  1.591431

In [100]: stacked.unstack(0)
Out[100]:
first      bar      baz
second
one  A -0.727965 -0.339355
     B -0.589346  0.593616
two  A  0.339969  0.884345
     B -0.693205  1.591431

```

Pivot tables

See the section on [Pivot Tables](#).

```

In [101]: df = pd.DataFrame(
.....:     {
.....:         "A": ["one", "one", "two", "three"] * 3,
.....:         "B": ["A", "B", "C"] * 4,
.....:         "C": ["foo", "foo", "foo", "bar", "bar", "bar"] * 2,
.....:         "D": np.random.randn(12),
.....:         "E": np.random.randn(12),
.....:     }
.....: )

In [102]: df
Out[102]:
   A  B  C      D      E
0  one A  foo -1.202872  0.047609
1  one B  foo -1.814470 -0.136473
2  two C  foo  1.018601 -0.561757
3  three A  bar -0.595447 -1.623033
4  one B  bar  1.395433  0.029399
5  one C  bar -0.392670 -0.542108
6  two A  foo  0.007207  0.282696
7  three B  foo  1.928123 -0.087302
8  one C  foo -0.055224 -1.575170
9  one A  bar  2.395985  1.771208
10 two B  bar  1.552825  0.816482
11 three C  bar  0.166599  1.100230

```

We can produce pivot tables from this data very easily:

```

In [103]: pd.pivot_table(df, values="D", index=["A", "B"], columns=["C"])
Out[103]:
           bar      foo
A  B
one A  2.395985 -1.202872
   B  1.395433 -1.814470
   C -0.392670 -0.055224
three A -0.595447      NaN
     B      NaN  1.928123
     C  0.166599      NaN
two  A      NaN  0.007207
     B  1.552825      NaN
     C      NaN  1.018601

```

Time series

pandas has simple, powerful, and efficient functionality for performing resampling operations during frequency conversion (e.g., converting secondly data into 5-minutely data). This is extremely common in, but not limited to, financial applications. See the [Time Series section](#).

```
In [104]: rng = pd.date_range("1/1/2012", periods=100, freq="S")
In [105]: ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)
In [106]: ts.resample("5Min").sum()
Out[106]:
2012-01-01    24182
Freq: 5T, dtype: int64
```

Time zone representation:

```
In [107]: rng = pd.date_range("3/6/2012 00:00", periods=5, freq="D")
In [108]: ts = pd.Series(np.random.randn(len(rng)), rng)
In [109]: ts
Out[109]:
2012-03-06    1.857704
2012-03-07   -1.193545
2012-03-08    0.677510
2012-03-09   -0.153931
2012-03-10    0.520091
Freq: D, dtype: float64
In [110]: ts_utc = ts.tz_localize("UTC")
In [111]: ts_utc
Out[111]:
2012-03-06 00:00:00+00:00    1.857704
2012-03-07 00:00:00+00:00   -1.193545
2012-03-08 00:00:00+00:00    0.677510
2012-03-09 00:00:00+00:00   -0.153931
2012-03-10 00:00:00+00:00    0.520091
Freq: D, dtype: float64
```

Converting to another time zone:

```
In [112]: ts_utc.tz_convert("US/Eastern")
Out[112]:
2012-03-05 19:00:00-05:00    1.857704
2012-03-06 19:00:00-05:00   -1.193545
2012-03-07 19:00:00-05:00    0.677510
2012-03-08 19:00:00-05:00   -0.153931
2012-03-09 19:00:00-05:00    0.520091
Freq: D, dtype: float64
```

Converting between time span representations:

```
In [113]: rng = pd.date_range("1/1/2012", periods=5, freq="M")
```

```
In [114]: ts = pd.Series(np.random.randn(len(rng)), index=rng)
```

```
In [115]: ts
```

```
Out[115]:
2012-01-31    -1.475051
2012-02-29     0.722570
2012-03-31    -0.322646
2012-04-30    -1.601631
2012-05-31     0.778033
Freq: M, dtype: float64
```

```
In [116]: ps = ts.to_period()
```

```
In [117]: ps
```

```
Out[117]:
2012-01    -1.475051
2012-02     0.722570
2012-03    -0.322646
2012-04    -1.601631
2012-05     0.778033
Freq: M, dtype: float64
```

```
In [118]: ps.to_timestamp()
```

```
Out[118]:
2012-01-01    -1.475051
2012-02-01     0.722570
2012-03-01    -0.322646
2012-04-01    -1.601631
2012-05-01     0.778033
Freq: MS, dtype: float64
```

Converting between period and timestamp enables some convenient arithmetic functions to be used. In the following example, we convert a quarterly frequency with year ending in November to 9am of the end of the month following the quarter end:

```
In [119]: prng = pd.period_range("1990Q1", "2000Q4", freq="Q-NOV")
```

```
In [120]: ts = pd.Series(np.random.randn(len(prng)), prng)
```

```
In [121]: ts.index = (prng.asfreq("M", "e") + 1).asfreq("H", "s") + 9
```

```
In [122]: ts.head()
```

```
Out[122]:
1990-03-01 09:00    -0.289342
1990-06-01 09:00     0.233141
1990-09-01 09:00    -0.223540
1990-12-01 09:00     0.542054
1991-03-01 09:00    -0.688585
Freq: H, dtype: float64
```

Categoricals

pandas can include categorical data in a `DataFrame`. For full docs, see the [categorical introduction](#) and the [API documentation](#).

```
In [123]: df = pd.DataFrame(
.....:     {"id": [1, 2, 3, 4, 5, 6], "raw_grade": ["a", "b", "b", "a", "a",
.....:     "e"]}
.....: )
.....:
```

Converting the raw grades to a categorical data type:


```
In [124]: df["grade"] = df["raw_grade"].astype("category")
```

```
In [125]: df["grade"]
```

```
Out[125]:
```

```
0    a
1    b
2    b
3    a
4    a
5    e
```

```
Name: grade, dtype: category
```

```
Categories (3, object): ['a', 'b', 'e']
```

Rename the categories to more meaningful names (assigning to `Series.cat.categories()` is in place!):

```
In [126]: df["grade"].cat.categories = ["very good", "good", "very bad"]
```

Reorder the categories and simultaneously add the missing categories (methods under `Series.cat()` return a new `Series` by default):

```
In [127]: df["grade"] = df["grade"].cat.set_categories(
.....:     ["very bad", "bad", "medium", "good", "very good"]
.....: )
.....:
```

```
In [128]: df["grade"]
```

```
Out[128]:
```

```
0    very good
1         good
2         good
3    very good
4    very good
5     very bad
```

```
Name: grade, dtype: category
```

```
Categories (5, object): ['very bad', 'bad', 'medium', 'good', 'very good']
```

Sorting is per order in the categories, not lexical order:

```
In [129]: df.sort_values(by="grade")
```

```
Out[129]:
```

	id	raw_grade	grade
5	6	e	very bad
1	2	b	good
2	3	b	good
0	1	a	very good
3	4	a	very good
4	5	a	very good

Grouping by a categorical column also shows empty categories:

```
In [130]: df.groupby("grade").size()
```

```
Out[130]:
```

```
grade
very bad    1
bad         0
medium      0
good        2
very good   3
dtype: int64
```

Plotting

See the [Plotting](#) docs.


We use the standard convention for referencing the matplotlib API:

```
In [131]: import matplotlib.pyplot as plt
```

```
In [132]: plt.close("all")
```

The `close()` method is used to `close` a figure window:

```
In [133]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000",
periods=1000))
In [134]: ts = ts.cumsum()
In [135]: ts.plot();
```



If running under Jupyter Notebook, the plot will appear on `plot()`. Otherwise use `matplotlib.pyplot.show` to show it or `matplotlib.pyplot.savefig` to write it to a file.

```
In [136]: plt.show();
```

On a DataFrame, the `plot()` method is a convenience to plot all of the columns with labels:

```
In [137]: df = pd.DataFrame(
.....:     np.random.randn(1000, 4), index=ts.index, columns=["A", "B", "C",
.....:     "D"]
.....: )
.....:
In [138]: df = df.cumsum()
In [139]: plt.figure();
In [140]: df.plot();
In [141]: plt.legend(loc='best');
```

🔍 Search the docs ...

10 minutes to pandas

[Intro to data structures](#)

[Essential basic functionality](#)

[IO tools \(text, CSV, HDF5, ...\)](#)

[Indexing and selecting data](#)

[MultiIndex / advanced indexing](#)

[Merge, join, concatenate and compare](#)

[Reshaping and pivot tables](#)


[Working with text data](#)

[Working with missing data](#)

[Duplicate Labels](#)

[Categorical data](#)

[Nullable integer data type](#)
[Nullable Boolean data type](#)
[Chart Visualization](#)
[Table Visualization](#)
[Computational tools](#)



Getting data in/out

CSV

[Writing to a csv file:](#)

```
In [142]: df.to_csv("foo.csv")
```

[Reading from a csv file:](#)

```
In [143]: pd.read_csv("foo.csv")
```

```
Out[143]:
```

	Unnamed: 0	A	B	C	D
0	2000-01-01	0.350262	0.843315	1.798556	0.782234
1	2000-01-02	-0.586873	0.034907	1.923792	-0.562651
2	2000-01-03	-1.245477	-0.963406	2.269575	-1.612566
3	2000-01-04	-0.252830	-0.498066	3.176886	-1.275581
4	2000-01-05	-1.044057	0.118042	2.768571	0.386039
...
995	2002-09-22	-48.017654	31.474551	69.146374	-47.541670
996	2002-09-23	-47.207912	32.627390	68.505254	-48.828331
997	2002-09-24	-48.907133	31.990402	67.310924	-49.391051
998	2002-09-25	-50.146062	33.716770	67.717434	-49.037577
999	2002-09-26	-49.724318	33.479952	68.108014	-48.822030

```
[1000 rows x 5 columns]
```

HDF5

Reading and writing to [HDFStores](#).

[Writing to a HDF5 Store:](#)

```
In [144]: df.to_hdf("foo.h5", "df")
```

Reading from a HDF5 Store:

```
In [145]: pd.read_hdf("foo.h5", "df")
Out[145]:
```

	A	B	C	D
2000-01-01	0.350262	0.843315	1.798556	0.782234
2000-01-02	-0.586873	0.034907	1.923792	-0.562651
2000-01-03	-1.245477	-0.963406	2.269575	-1.612566
2000-01-04	-0.252830	-0.498066	3.176886	-1.275581
2000-01-05	-1.044057	0.118042	2.768571	0.386039
...
2002-09-22	-48.017654	31.474551	69.146374	-47.541670
2002-09-23	-47.207912	32.627390	68.505254	-48.828331
2002-09-24	-48.907133	31.990402	67.310924	-49.391051
2002-09-25	-50.146062	33.716770	67.717434	-49.037577
2002-09-26	-49.724318	33.479952	68.108014	-48.822030

```
[1000 rows x 4 columns]
```

Excel

Reading and writing to [MS Excel](#).

Writing to an excel file:

```
In [146]: df.to_excel("foo.xlsx", sheet_name="Sheet1")
```

Reading from an excel file:

```
In [147]: pd.read_excel("foo.xlsx", "Sheet1", index_col=None, na_values=["NA"])
Out[147]:
```

Unnamed: 0	A	B	C	D	
0	2000-01-01	0.350262	0.843315	1.798556	0.782234
1	2000-01-02	-0.586873	0.034907	1.923792	-0.562651
2	2000-01-03	-1.245477	-0.963406	2.269575	-1.612566
3	2000-01-04	-0.252830	-0.498066	3.176886	-1.275581
4	2000-01-05	-1.044057	0.118042	2.768571	0.386039
...
995	2002-09-22	-48.017654	31.474551	69.146374	-47.541670
996	2002-09-23	-47.207912	32.627390	68.505254	-48.828331
997	2002-09-24	-48.907133	31.990402	67.310924	-49.391051
998	2002-09-25	-50.146062	33.716770	67.717434	-49.037577
999	2002-09-26	-49.724318	33.479952	68.108014	-48.822030

```
[1000 rows x 5 columns]
```

Gotchas

If you are attempting to perform an operation you might see an exception like:

```
>>> if pd.Series([False, True, False]):
...     print("I was true")
Traceback
...
ValueError: The truth value of an array is ambiguous. Use a.empty(), a.any() or
a.all().
```

See [Comparisons](#) for an explanation and what to do.

See [Gotchas](#) as well.

◀ Previous
[User Guide](#)

[Intro to data structures](#) ▶