

# ML FSN Programming

## ▼ Linear Regression (CO2 Emissions )

```
1 import seaborn as sns
2 import pandas as pd
3 import numpy as np
```


### Downloading Data

```
1 !wget -O FuelConsumption.csv https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/CF-CO2-Emissions-by-Country.csv

--2022-06-20 08:54:07-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/CF-CO2-Emissions-by-Country.csv
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud): 52.91.142.100
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud): 200 OK
HTTP request sent, awaiting response... 200 OK
Length: 72629 (71K) [text/csv]
Saving to: 'FuelConsumption.csv'

FuelConsumption.csv 100%[=====>] 70.93K --.-KB/s in 0.04s

2022-06-20 08:54:07 (1.93 MB/s) - 'FuelConsumption.csv' saved [72629/72629]
```



### Reading the data in

```
1 df = pd.read_csv("FuelConsumption.csv")
2
3 # take a look at the dataset
4 df.head()
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUELCONSUMPTION_CITY
0	2014	ACURA	ILX	COMPACT	2.0	4	AS5	

```
1 df.corr()
```

	MODELYEAR	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY
MODELYEAR	NaN	NaN	NaN	NaN
ENGINE SIZE	NaN	1.000000	0.934011	0.832225
CYLINDERS	NaN	0.934011	1.000000	0.796473
FUELCONSUMPTION_CITY	NaN	0.832225	0.796473	1.000000
FUELCONSUMPTION_HWY	NaN	0.778746	0.724594	0.965700
FUELCONSUMPTION_COMB	NaN	0.819482	0.776788	0.995500
FUELCONSUMPTION_COMB_MPG	NaN	-0.808554	-0.770430	-0.935600
CO2EMISSIONS	NaN	0.874154	0.849685	0.898039



```
1 df.corr()['CO2EMISSIONS'].sort_values(ascending=False)
```

```
CO2EMISSIONS      1.000000
FUELCONSUMPTION_CITY  0.898039
FUELCONSUMPTION_COMB  0.892129
ENGINE SIZE        0.874154
FUELCONSUMPTION_HWY  0.861748
CYLINDERS          0.849685
FUELCONSUMPTION_COMB_MPG -0.906394
MODELYEAR          NaN
Name: CO2EMISSIONS, dtype: float64
```

Let's select some features that we want to use for regression.


```
1 cdf = df[['ENGINE SIZE', 'CYLINDERS', 'FUELCONSUMPTION_CITY', 'FUELCONSUMPTION_HWY', 'FUELCONSUMPTION_COMB_MPG']]
2 cdf.head(9)
```

	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY	FUELCONSUMPTION_COMB
0	2.0	4	9.9	6.7	8.1
1	2.4	4	11.2	7.7	9.4
2	1.5	4	6.0	5.8	5.9

Let's plot Emission values with respect to Engine size:

```
1 import seaborn as sns
2 sns.scatterplot(data=cdf,x=cdf.ENGINE SIZE,y=cdf.CO2EMISSIONS,hue=cdf.CYLINDERS)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff42d6a6a90>



```
1 cdf.columns
```

```
Index(['ENGINE SIZE', 'CYLINDERS', 'FUELCONSUMPTION_CITY',
      'FUELCONSUMPTION_HWY', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS'],
      dtype='object')
```

```
1 cdf_x=cdf[['ENGINE SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']]
```

```
1 cdf_y=cdf[['CO2EMISSIONS']]
```

```
1 from sklearn.model_selection import train_test_split
2 X_train,X_test,y_train,y_test=train_test_split(cdf_x,cdf_y,test_size=0.2,random_state=4)
```

```
1 X_train.head()
```


ENGINE SIZE   CYLINDERS   FUEL CONSUMPTION\_COMB

```
1 y_train.head()
```

	CO2EMISSIONS
409	196
773	205
146	264
776	246
381	336

```
1 sns.scatterplot(X_train.ENGINE SIZE,y_train.CO2EMISSIONS)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass   
 FutureWarning  
<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff42b982d90>



## Multiple Regression Model

```
1 from sklearn import linear_model
2 regr = linear_model.LinearRegression()
3 regr
```

LinearRegression()

```
1 regr.fit (X_train, y_train)
```

LinearRegression()

```
1 regr.intercept_
```

```
array([65.2578757])
```

```
1 regr.coef_
```

```
array([[10.24537129,  7.64355532,  9.68132732]])
```

```
1
```

## Prediction

```
1 y_hat= regr.predict(X_test)
```

```
2 y_hat[0:5]
```

```
array([[259.39421287],
       [216.04051098],
       [255.40887315],
       [261.21766954],
       [294.43880893]])
```

```
1 x = X_test
```

```
2 y = y_test
```

```
1 np.mean(y_hat-y)
```

```
CO2EMISSIONS    -1.375796
dtype: float64
```

```
1 #Residual sum of squares:
```

```
2 np.mean((y_hat - y) ** 2)
```

```
CO2EMISSIONS    408.37553
dtype: float64
```

```
1 # Explained variance score: 1 is perfect prediction
```

```
2 #print('Variance score: %.2f' %
```

```
3 regr.score(x, y)
```

```
0.890023090970219
```

## ▼ KNN (Customer Category)

```
1 !wget -O teleCust1000t.csv https://cf-courses-data.s3.us.cloud-object-storage.appdomain
```

```
--2022-06-20 09:03:33--  https://cf-courses-data.s3.us.cloud-object-storage.appdomain
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud): 34.100.100.100
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud): 34.100.100.100
```

HTTP request sent, awaiting response... 200 OK  
 Length: 36047 (35K) [text/csv]  
 Saving to: 'teleCust1000t.csv'

teleCust1000t.csv 100%[=====>] 35.20K --.-KB/s in 0.02s

2022-06-20 09:03:33 (1.39 MB/s) - 'teleCust1000t.csv' saved [36047/36047]

```
1 df = pd.read_csv('teleCust1000t.csv')
2 df.head()
```

	region	tenure	age	marital	address	income	ed	employ	retire	gender	reside
0	2	13	44	1	9	64.0	4	5	0.0	0	2
1	3	11	33	1	7	136.0	5	5	0.0	0	6
2	3	68	52	1	24	116.0	1	29	0.0	1	2
3	2	33	33	0	12	33.0	2	0	0.0	1	1
4	2	23	30	1	9	30.0	1	2	0.0	0	4

```
1 df.custcat.value_counts()
```

```
3    281
1    266
4    236
2    217
Name: custcat, dtype: int64
```

```
1 df.corr()['custcat'].sort_values(ascending=False)
```

```
custcat    1.000000
ed          0.193864
tenure      0.166691
income      0.134525
employ      0.110011
marital     0.083836
reside      0.082022
address     0.067913
age         0.056909
retire      0.008908
gender     -0.004966
region     -0.023771
Name: custcat, dtype: float64
```

```
1 df.columns
```

```
Index(['region', 'tenure', 'age', 'marital', 'address', 'income', 'ed',
      'employ', 'retire', 'gender', 'reside', 'custcat'],
      dtype='object')
```

Normalize Data Data Standardization gives the data zero mean and unit variance, it is good practice, especially for algorithms such as KNN which is based on the distance of data points:

### TO use scikit learn convert pandas df to np array

```
1 X = df[['region', 'tenure', 'age', 'marital', 'address', 'income', 'ed', 'employ', 'retire', 'gender', 'reside']]
2 X[0:5]
```

	region	tenure	age	marital	address	income	ed	employ	retire	gender	reside
0	2	13	44	1	9	64.0	4	5	0.0	0	2
1	3	11	33	1	7	136.0	5	5	0.0	0	6
2	3	68	52	1	24	116.0	1	29	0.0	1	2
3	2	33	33	0	12	33.0	2	0	0.0	1	1
4	2	23	30	1	9	30.0	1	2	0.0	0	4

its still df, so use .values

```
1 X = df[['region', 'tenure', 'age', 'marital', 'address', 'income', 'ed', 'employ', 'retire', 'gender', 'reside']].values
2 X[0:5]
```

```
array([[ 2., 13., 44.,  1.,  9., 64.,  4.,  5.,  0.,  0.,  2.],
       [ 3., 11., 33.,  1.,  7., 136.,  5.,  5.,  0.,  0.,  6.],
       [ 3., 68., 52.,  1., 24., 116.,  1., 29.,  0.,  1.,  2.],
       [ 2., 33., 33.,  0., 12.,  33.,  2.,  0.,  0.,  1.,  1.],
       [ 2., 23., 30.,  1.,  9.,  30.,  1.,  2.,  0.,  0.,  4.]])
```

```
1 Y = df[['custcat']].values
2 Y[0:5]
```

```
array([[1],
       [4],
       [3],
       [1],
       [3]])
```

```
1 X.shape
```

```
(1000, 11)
```

```
1 Y.shape
```

```
(1000, 1)
```

```
1 from sklearn import preprocessing
```

```
1 X = preprocessing.StandardScaler().fit(X).transform(X.astype(float))
2 X[0:5]
```

```
array([[ -0.02696767, -1.055125  ,  0.18450456,  1.0100505 , -0.25303431,
        -0.12650641,  1.0877526 , -0.5941226 , -0.22207644, -1.03459817,
        -0.23065004],
       [ 1.19883553, -1.14880563, -0.69181243,  1.0100505 , -0.4514148 ,
        0.54644972,  1.9062271 , -0.5941226 , -0.22207644, -1.03459817,
        2.55666158],
       [ 1.19883553,  1.52109247,  0.82182601,  1.0100505 ,  1.23481934,
        0.35951747, -1.36767088,  1.78752803, -0.22207644,  0.96655883,
        -0.23065004],
       [-0.02696767, -0.11831864, -0.69181243, -0.9900495 ,  0.04453642,
        -0.41625141, -0.54919639, -1.09029981, -0.22207644,  0.96655883,
        -0.92747794],
       [-0.02696767, -0.58672182, -0.93080797,  1.0100505 , -0.25303431,
        -0.44429125, -1.36767088, -0.89182893, -0.22207644, -1.03459817,
        1.16300577]])
```

```
1 X.shape,Y.shape
```

```
((1000, 11), (1000, 1))
```

```
1 from sklearn.model_selection import train_test_split
```

```
1 X_train, Y_train , X_test , Y_test = train_test_split(X,Y,test_size=0.2,random_state=4)
2 #order is wrong , X train then X_test
```

```
1 X_train.shape , Y_train.shape
```

```
((800, 11), (200, 11))
```

```
1 X_train, X_test, Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=4)
```

```
1 X_train.shape , Y_train.shape
```

```
((800, 11), (800, 1))
```

```
1 from sklearn.neighbors import KNeighborsClassifier
```

```
1 k=4
2 KNN = KNeighborsClassifier(n_neighbors = k).fit(X_train,Y_train)
3 KNN
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
KNeighborsClassifier(n_neighbors=4)
```



```
1 Y_hat = KNN.predict(X_test)

1 Y_hat[0:5]

array([1, 1, 3, 2, 4])
```

for multilabel classifier , we use jaccard score

```
1 from sklearn import metrics
```

```
1 metrics.accuracy_score(Y_test,Y_hat)

0.32
```

```
1 metrics.accuracy_score(Y_hat,Y_test)


0.32
```

```
1 k_max = 11
2 k_score_l=[]
3 for i in range(1,k_max):
4     KNN = KNeighborsClassifier(n_neighbors=i).fit(X_train,Y_train)
5     Y_hat = KNN.predict(X_test)
6     scr = metrics.accuracy_score(Y_test,Y_hat)
7     k_score_l.append(scr)
8 k_score_l
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
[0.3, 0.29, 0.315, 0.32, 0.315, 0.31, 0.335, 0.325, 0.34, 0.33]
```

```
1 sns.lineplot(y=np.array(k_score_l),x=np.arange(1,11))
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7ff42178d090>



1 # k=9 is max score , do that score


## ▼ Decision Tree ( Drug A B C)

```
1 !wget -O drug200.csv https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud
```

```
--2022-06-20 11:26:40-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)
HTTP request sent, awaiting response... 200 OK
Length: 5827 (5.7K) [text/csv]
Saving to: 'drug200.csv'
```

```
drug200.csv      100%[=====>]    5.69K  --.-KB/s    in 0s
```

```
2022-06-20 11:26:40 (1.16 GB/s) - 'drug200.csv' saved [5827/5827]
```



```
1 from sklearn.tree import DecisionTreeClassifier
```

```
1 drug = pd.read_csv("drug200.csv")
2 drug
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

```
1 drug.head()
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

```
1 X = drug[['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K']].values
2 X[0:5]
```

```
array([[23, 'F', 'HIGH', 'HIGH', 25.355],
       [47, 'M', 'LOW', 'HIGH', 13.093],
       [47, 'M', 'LOW', 'HIGH', 10.114],
       [28, 'F', 'NORMAL', 'HIGH', 7.798],
       [61, 'F', 'LOW', 'HIGH', 18.043]], dtype=object)
```

```
1 Y = drug[['Drug']].values
2 Y[0:5]
```

```
array(['drugY',
       'drugC',
       'drugC',
       'drugX',
       'drugY'], dtype=object)
```

```
1 drug.columns
```

```
Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

```
1 from sklearn import preprocessing
2
3
4 LE_Sex = preprocessing.LabelEncoder()
5 LE_Sex.fit(['M', 'F'])
6 X[:,1]= LE_Sex.transform(X[:,1])
7 X[:5]
```

```
array([[23, 0, 'HIGH', 'HIGH', 25.355],
       [47, 1, 'LOW', 'HIGH', 13.093],
       [47, 1, 'LOW', 'HIGH', 10.114],
       [28, 0, 'NORMAL', 'HIGH', 7.798],
       [61, 0, 'LOW', 'HIGH', 18.043]], dtype=object)
```

```
1 drug.BP.unique()
```

```
array(['HIGH', 'LOW', 'NORMAL'], dtype=object)
```

```
1 LE_BP = preprocessing.LabelEncoder()
2 LE_BP.fit(['HIGH', 'LOW', 'NORMAL'])
3 X[:,2]=LE_BP.transform(X[:,2])
4 X[:5]
```

```
array([[23, 0, 0, 'HIGH', 25.355],
       [47, 1, 1, 'HIGH', 13.093],
       [47, 1, 1, 'HIGH', 10.114],
       [28, 0, 2, 'HIGH', 7.798],
       [61, 0, 1, 'HIGH', 18.043]], dtype=object)
```

```
1 drug.Cholesterol.unique()
```

```
array(['HIGH', 'NORMAL'], dtype=object)
```

```
1 LE_Chol = preprocessing.LabelEncoder()
2 LE_Chol.fit(['NORMAL', 'HIGH'])
3 X[:,3] = LE_Chol.transform(X[:,3])
```

```
1 X
```

```
array([[23, 0, 0, 0, 25.355],
       [47, 1, 1, 0, 13.093],
       [47, 1, 1, 0, 10.114],
       [28, 0, 2, 0, 7.798],
       [61, 0, 1, 0, 18.043],
       [22, 0, 2, 0, 8.607],
       [49, 0, 2, 0, 16.275],
       [41, 1, 1, 0, 11.037],
       [60, 1, 2, 0, 15.171],
       [43, 1, 1, 1, 19.368],
       [47, 0, 1, 0, 11.767],
       [34, 0, 0, 1, 19.199],
       [43, 1, 1, 0, 15.376],
       [74, 0, 1, 0, 20.942],
       [50, 0, 2, 0, 12.703],
       [16, 0, 0, 1, 15.516],
       [69, 1, 1, 1, 11.455],
       [43, 1, 0, 0, 13.972],
       [23, 1, 1, 0, 7.298],
       [32, 0, 0, 1, 25.974],
       [57, 1, 1, 1, 19.128],
       [63, 1, 2, 0, 25.917],
       [47, 1, 1, 1, 30.568],
```

```
[48, 0, 1, 0, 15.036],
[33, 0, 1, 0, 33.486],
[28, 0, 0, 1, 18.809],
[31, 1, 0, 0, 30.366],
[49, 0, 2, 1, 9.381],
[39, 0, 1, 1, 22.697],
[45, 1, 1, 0, 17.951],
[18, 0, 2, 1, 8.75],
[74, 1, 0, 0, 9.567],
[49, 1, 1, 1, 11.014],
[65, 0, 0, 1, 31.876],
[53, 1, 2, 0, 14.133],
[46, 1, 2, 1, 7.285],
[32, 1, 0, 1, 9.445],
[39, 1, 1, 1, 13.938],
[39, 0, 2, 1, 9.709],
[15, 1, 2, 0, 9.084],
[73, 0, 2, 0, 19.221],
[58, 0, 0, 1, 14.239],
[50, 1, 2, 1, 15.79],
[23, 1, 2, 0, 12.26],
[50, 0, 2, 1, 12.295],
[66, 0, 2, 1, 8.107],
[37, 0, 0, 0, 13.091],
[68, 1, 1, 0, 10.291],
[23, 1, 2, 0, 31.686],
[28, 0, 1, 0, 19.796],
[58, 0, 0, 0, 19.416],
[67, 1, 2, 1, 10.898],
[62, 1, 1, 1, 27.183],
[24, 0, 0, 1, 18.457],
[68, 0, 0, 1, 10.189],
[26, 0, 1, 0, 14.16],
[65, 1, 0, 1, 11.34],
_ _ _ _ _
```

```
1 from sklearn import preprocessing
2 from sklearn.model_selection import train_test_split
```

```
1 X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=4)
```

```
1 X_train.shape
```

```
(160, 5)
```

```
1 X_test.shape
```

```
(40, 5)
```

```
1 DecTree = DecisionTreeClassifier(criterion='entropy',max_depth=4)
2 DecTree
```

```
DecisionTreeClassifier(criterion='entropy', max_depth=4)
```

```
1 DecTree.fit(X_train,Y_train)
```

```
DecisionTreeClassifier(criterion='entropy', max_depth=4)
```

```
1 Y_hat=DecTree.predict(X_test)
2 Y_hat
```

```
array(['drugY', 'drugY', 'drugY', 'drugY', 'drugC', 'drugX', 'drugY',
       'drugY', 'drugY', 'drugA', 'drugA', 'drugX', 'drugA', 'drugY',
       'drugY', 'drugY', 'drugY', 'drugX', 'drugA', 'drugC', 'drugX',
       'drugC', 'drugA', 'drugX', 'drugC', 'drugB', 'drugX', 'drugY',
       'drugX', 'drugY', 'drugB', 'drugC', 'drugX', 'drugX', 'drugY',
       'drugY', 'drugA', 'drugA', 'drugX', 'drugY'], dtype=object)
```

```
1 from sklearn import metrics
2 metrics.accuracy_score(Y_hat,Y_test)
```

```
0.95
```


## ▼ SVM (Cancer Data Set )

```
1 !wget -O cell_samples.csv https://cf-courses-data.s3.us.cloud-object-storage.appdomain
```

```
--2022-06-20 16:30:41-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-
HTTP request sent, awaiting response... 200 OK
Length: 19975 (20K) [text/csv]
Saving to: 'cell_samples.csv'
```

```
cell_samples.csv 100%[=====>] 19.51K --.-KB/s in 0s
```

```
2022-06-20 16:30:41 (165 MB/s) - 'cell_samples.csv' saved [19975/19975]
```



```
1 import pandas as pd
```

```
1 df = pd.read_csv("cell_samples.csv")
2 df.head()
```

	ID	Clump	UnifSize	UnifShape	MargAdh	SingEpiSize	BareNuc	BlandChrom	Nc
0	1000025	5	1	1	1	2	1	3	
1	1002945	5	4	4	5	7	10	3	
2	1015425	3	1	1	1	2	2	3	
3	1016277	6	8	8	1	3	4	3	
4	1017023	4	1	1	3	2	1	3	



```
1 df.columns
```

```
Index(['ID', 'Clump', 'UnifSize', 'UnifShape', 'MargAdh', 'SingEpiSize',
      'BareNuc', 'BlandChrom', 'NormNucl', 'Mit', 'Class'],
      dtype='object')
```

```
1 df=df[['Clump', 'UnifSize', 'UnifShape', 'MargAdh', 'SingEpiSize', 'BareNuc', 'BlandChr
```

```
1 df
```

	Clump	UnifSize	UnifShape	MargAdh	SingEpiSize	BareNuc	BlandChrom	NormNucl
0	5	1	1	1	2	1	3	1
1	5	4	4	5	7	10	3	2
2	3	1	1	1	2	2	3	1
3	6	8	8	1	3	4	3	7
4	4	1	1	3	2	1	3	1
...	...	...	...	...	...	...	...	...
694	3	1	1	1	3	2	1	1
695	2	1	1	1	2	1	1	1
696	5	10	10	3	7	3	8	10
697	4	8	6	4	3	4	10	6
698	4	8	8	5	4	5	10	4

699 rows × 10 columns



```
1 df.columns
```

```
Index(['Clump', 'UnifSize', 'UnifShape', 'MargAdh', 'SingEpiSize', 'BareNuc',
      'BlandChrom', 'NormNucl', 'Mit', 'Class'],
      dtype='object')
```

```
1 import seaborn as sns
2 sns.scatterplot(df.Clump, df.UnifSize, hue=df.Class)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f34823a3790>
```



```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 699 entries, 0 to 698
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Clump       699 non-null    int64
1   UnifSize    699 non-null    int64
2   UnifShape   699 non-null    int64
3   MargAdh     699 non-null    int64
4   SingEpiSize 699 non-null    int64
5   BareNuc     699 non-null    object
6   BlandChrom  699 non-null    int64
7   NormNuc1    699 non-null    int64
8   Mit         699 non-null    int64
9   Class       699 non-null    int64
dtypes: int64(9), object(1)
memory usage: 54.7+ KB
```

```
1 df['BareNuc'] = df['BareNuc'].astype('int')
```

```
ValueError                                Traceback (most recent call last)
```

```
<ipython-input-16-f54d3735dc60> in <module>()
```

```
----> 1 df['BareNuc'] = df['BareNuc'].astype('int')
```

7 frames

```
/usr/local/lib/python3.7/dist-packages/pandas/_libs/lib.pyx in
pandas._libs.lib.astype_intsafe()
```

```
ValueError: invalid literal for int() with base 10: '?'
```

SEARCH STACK OVERFLOW

```
1 df = df[pd.to_numeric(df['BareNuc'], errors='coerce').notnull()]
2 df['BareNuc'] = df['BareNuc'].astype('int')
3 df.dtypes
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/using.html>

```
Clump      int64
UnifSize    int64
UnifShape   int64
MargAdh     int64
```



```

SingEpiSize    int64
BareNuc        int64
BlandChrom     int64
NormNucl       int64
Mit            int64
Class          int64
dtype: object

```

```
1 df.head()
```

	Clump	UnifSize	UnifShape	MargAdh	SingEpiSize	BareNuc	BlandChrom	NormNucl	Mit
0	5	1	1	1	2	1	3	1	1
1	5	4	4	5	7	10	3	2	1
2	3	1	1	1	2	2	3	1	1
3	6	8	8	1	3	4	3	7	1
4	4	1	1	3	2	1	3	1	1

```

1 X=df.values[:, :-1]
2 X

```

```

array([[ 5,  1,  1, ...,  3,  1,  1],
       [ 5,  4,  4, ...,  3,  2,  1],
       [ 3,  1,  1, ...,  3,  1,  1],
       ...,
       [ 5, 10, 10, ...,  8, 10,  2],
       [ 4,  8,  6, ..., 10,  6,  1],
       [ 4,  8,  8, ..., 10,  4,  1]])

```

```

1 Y=df.values[:, -1]
2 Y

```

```

array([2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 4, 2, 4, 4, 2, 2, 4, 2, 4, 4,
       2, 2, 4, 2, 2, 2, 2, 2, 4, 2, 2, 2, 4, 2, 4, 4, 4, 4, 4, 4, 2, 4, 4, 2, 4, 2, 4,
       4, 2, 2, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 2, 4, 4, 2, 4, 2, 4,
       4, 2, 2, 4, 2, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 4, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 4, 4, 4, 2, 4, 4, 4, 4, 2, 4, 2, 4,
       4, 4, 2, 2, 2, 4, 2, 2, 2, 4, 4, 4, 2, 4, 2, 4, 2, 2, 2, 4, 2, 2, 4, 2,
       2, 2, 2, 2, 2, 2, 4, 4, 2, 2, 2, 2, 2, 4, 2, 2, 4, 2, 4, 2, 2,
       4, 4, 2, 2, 2, 4, 4, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 4, 2, 4, 2, 2,
       2, 4, 4, 2, 4, 4, 4, 2, 4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 4, 4, 2, 2,
       2, 4, 4, 2, 2, 2, 4, 4, 2, 4, 4, 4, 2, 2, 4, 4, 4, 4, 4, 2,
       4, 4, 2, 4, 4, 4, 2, 4, 2, 4, 4, 4, 4, 2, 2, 2, 2, 2, 2, 4, 4, 2,
       2, 4, 4, 2, 4, 2, 4, 4, 2, 2, 4, 2, 2, 2, 4, 2, 2, 4, 4, 2, 2, 4,
       2, 4, 2, 2, 4, 2, 4, 4, 4, 2, 2, 4, 4, 2, 4, 2, 2, 4, 4, 2, 2, 2,
       4, 2, 2, 2, 4, 4, 2, 2, 2, 4, 2, 2, 4, 4, 4, 4, 4, 2, 2, 2, 2,
       4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 4, 2,
       2, 2, 2, 4, 2, 4, 2, 2, 2, 2, 4, 2, 2, 2, 4, 2, 4, 2, 2, 2, 2,
       2, 2, 4, 4, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 4, 2,

```

```
4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 4, 2, 2, 4, 4, 2, 2, 2, 4, 4, 4, 2, 4, 2, 4, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 4, 4, 2, 2, 2, 4, 2,
2, 4, 4, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
4, 2, 2, 4, 4, 4, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 4, 4, 2, 2, 2, 4,
2, 4, 2, 4, 4, 4, 2, 4, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 2, 4,
2, 4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2,
2, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2,
2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 2, 2, 2, 2,
2, 2, 2, 2, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 4, 4,
4])
```

```
1 X.shape,Y.shape
```

```
((683, 9), (683,))
```

```
1 from sklearn.model_selection import train_test_split
2 X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=4)
```

SVM algorithm offers a choice of kernel functions for performing its processing.

mapping data into a higher dimensional space is called kernelling.

The mathematical function used for the transformation is known as the kernel function, and can be of different types, such as:

1.Linear 2.Polynomial 3.Radial basis function (RBF) 4.Sigmoid

```
1 from sklearn import svm
```

```
1 SVM = svm.SVC(kernel='rbf')
2 SVM.fit(X_train,Y_train)
```

```
SVC()
```

```
1 Y_hat=SVM.predict(X_test)
2 Y_hat
```

```
array([2, 4, 2, 4, 2, 2, 2, 2, 4, 2, 2, 4, 4, 4, 4, 2, 2, 2, 2, 2, 4, 2,
4, 4, 4, 4, 2, 2, 4, 4, 4, 2, 4, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 4,
4, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 4, 4, 2, 4, 4,
4, 2, 2, 2, 4, 4, 2, 2, 2, 4, 2, 2, 4, 4, 2, 2, 2, 2, 4, 4, 2, 4,
2, 2, 4, 4, 2, 2, 2, 4, 2, 2, 2, 4, 2, 4, 2, 2, 4, 2, 4, 2, 2, 4,
2, 2, 4, 2, 2, 2, 2, 2, 4, 4, 4, 4, 4, 2, 2, 4, 2, 2, 4, 2, 4, 2,
2, 2, 2, 2, 4])
```

```
1 from sklearn import metrics
2 metrics.accuracy_score(Y_test,Y_hat)
```

```
0.9635036496350365
```

```
1 from sklearn.metrics import confusion_matrix
```

```
1 cnf_matrix = confusion_matrix(Y_test, Y_hat, labels=[2,4])
2 cnf_matrix
```

```
array([[85,  5],
       [ 0, 47]])
```

## ▼ K MEANS clustering (Customer Segmentation )

```
1 !wget -O Cust_Segmentation.csv https://cf-courses-data.s3.us.cloud-object-storage.app
```

```
→ --2022-06-20 10:42:27-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud):
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud):
HTTP request sent, awaiting response... 200 OK
Length: 33426 (33K) [text/csv]
Saving to: 'Cust_Segmentation.csv'
```

```
Cust_Segmentation.c 100%[=====] 32.64K --.-KB/s in 0.02s
```

```
2022-06-20 10:42:27 (1.30 MB/s) - 'Cust_Segmentation.csv' saved [33426/33426]
```

```
1 df = pd.read_csv("Cust_Segmentation.csv")
2 df.head()
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	Defaulted	Address	DebtIncome
0	1	41	2	6	19	0.124	1.073	0.0	NBA001	
1	2	47	1	26	100	4.582	8.218	0.0	NBA021	
2	3	33	2	10	57	6.111	5.802	1.0	NBA013	
3	4	29	2	4	19	0.681	0.516	0.0	NBA009	

As you can see, Address in this dataset is a categorical variable. The k-means algorithm isn't directly applicable to categorical variables because the Euclidean distance function isn't really meaningful for discrete variables. So, let's drop this feature and run clustering

```
1 df=df.drop('Address',axis=1)
2 df.head()
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	Defaulted	DebtIncomeRatio
0	1	41	2	6	19	0.124	1.073	0.0	6.3
1	2	47	1	26	100	4.582	8.218	0.0	12.8
2	3	33	2	10	57	6.111	5.802	1.0	20.9

```
1 from sklearn import preprocessing
```

```
1 from sklearn.preprocessing import StandardScaler
```

```
1 X = df.values[:,1:] #removing id column
2 X = np.nan_to_num(X)
3 Clus_dataSet = StandardScaler().fit_transform(X)
4 Clus_dataSet
```

```
array([[ 0.74291541,  0.31212243, -0.37878978, ..., -0.59048916,
        -0.52379654, -0.57652509],
       [ 1.48949049, -0.76634938,  2.5737211 , ...,  1.51296181,
        -0.52379654,  0.39138677],
       [-0.25251804,  0.31212243,  0.2117124 , ...,  0.80170393,
        1.90913822,  1.59755385],
       ...,
       [-1.24795149,  2.46906604, -1.26454304, ...,  0.03863257,
        1.90913822,  3.45892281],
       [-0.37694723, -0.76634938,  0.50696349, ..., -0.70147601,
        -0.52379654, -1.08281745],
       [ 2.1116364 , -0.76634938,  1.09746566, ...,  0.16463355,
        -0.52379654, -0.2340332 ]])
```

```
1 from sklearn.cluster import KMeans
```

```
1 k = 3
2 k_means = KMeans(init = "k-means++", n_clusters = k, n_init = 12)
3 k_means.fit(X)
```

```
KMeans(n_clusters=3, n_init=12)
```

```
1 labels = k_means.labels_
```

```
1 labels
```

```
array([1, 2, 1, 1, 0, 2, 1, 2, 1, 2, 2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1,
       1, 2, 2, 2, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 0,
       1, 2, 1, 1, 1, 2, 2, 1, 1, 2, 2, 1, 1, 1, 2, 1, 2, 1, 2, 2, 1, 1,
       2, 1, 1, 1, 2, 2, 2, 1, 1, 1, 1, 1, 2, 1, 2, 2, 0, 1, 1, 1, 1, 1,
       1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, 2,
       1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2,
       1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 1,
       2, 1, 2, 2, 1, 2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1,
       1, 1, 1, 2, 1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1, 2, 1, 2, 1, 1, 2,
       0, 1, 2, 1, 1, 1, 1, 1, 1, 0, 2, 1, 1, 1, 1, 2, 1, 1, 2, 2, 1, 2,
```

```
1, 2, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 0, 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1,
2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 2, 1,
2, 1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 1, 1, 1, 2, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 2, 1,
2, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 2, 1, 1, 2, 1,
1, 1, 1, 1, 2, 1, 1, 1, 0, 1, 1, 1, 2, 1, 2, 2, 2, 1, 1, 1, 2, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 2, 1, 1, 1, 1, 2, 1,
1, 1, 1, 2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 2, 1,
1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 2, 1,
1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 0,
1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1,
1, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 2, 1, 2, 1, 1, 2,
1, 2, 1, 1, 0, 1, 2, 1, 2, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 2, 1,
1, 1, 2, 2, 1, 1, 2, 1, 1, 1, 2, 1, 0, 1, 1, 2, 1, 1, 1, 1, 1, 1,
1, 2, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1,
1, 2, 1, 1, 2, 1, 2, 1, 2, 2, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 2,
1, 1, 1, 1, 2, 2, 1, 1, 2, 2, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 2, 1, 2, 1, 2, 2, 1, 1, 2,
1, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 0,
2, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2], dtype=int32)
```

```
1 df.head()
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	Defaulted	DebtIncomeRatio
0	1	41	2	6	19	0.124	1.073	0.0	6.3
1	2	47	1	26	100	4.582	8.218	0.0	12.8
2	3	33	2	10	57	6.111	5.802	1.0	20.9
3	4	29	2	4	19	0.681	0.516	0.0	6.3
4	5	47	1	31	253	9.308	8.908	0.0	7.2

```
1 df['Clus_km']=labels
2 df.head()
```

```
1 df.groupby('Clus_km').mean()
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	D
Clus_km								
0	410.166667	45.388889	2.666667	19.555556	227.166667	5.678444	10.907167	
1	432.468413	32.964561	1.614792	6.374422	31.164869	1.032541	2.104133	

```
1 df.groupby('Clus_km').mean().corr()
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	D
Customer Id	1.000000	-0.836471	-0.559038	-0.836649	-0.507732	-0.667231	-0.641298	
Age	-0.836471	1.000000	0.921998	1.000000	0.896823	0.966306	0.956912	
Edu	-0.559038	0.921998	1.000000	0.921872	0.998160	0.990595	0.994704	
Years Employed	-0.836649	1.000000	0.921872	1.000000	0.896679	0.966222	0.956817	
Income	-0.507732	0.896823	0.998160	0.896679	1.000000	0.980475	0.986641	
Card Debt	-0.667231	0.966306	0.990595	0.966222	0.980475	1.000000	0.999412	
Other Debt	-0.641298	0.956912	0.994704	0.956817	0.986641	0.999412	1.000000	
Defaulted	0.698847	-0.192589	0.202379	-0.192908	0.261389	0.066477	0.100629	

```
1 sns.swarmplot(x=df.Age,y=df.Income,hue=df.Clus_km)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 50.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 66.7
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 78.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 81.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 86.7
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 80.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 81.8
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 86.8
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 82.4
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 88.1
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 70.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 83.9
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 81.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 82.5
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 72.7
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 71.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 75.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 59.4
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 61.1
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 45.5
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 33.3
warnings.warn(msg, UserWarning)
```

```
1 sns.swarmplot(x=df.Edu,y=df.Income,hue=df.Clus_km)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 77.6
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 59.1
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 17.8
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 26.5
warnings.warn(msg, UserWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7ff420a2c9d0>
```

1

